

Concrete Strength Data EDA

Concrete, the ubiquitous building block, is the foundation of countless structures. Its strength is paramount, ensuring safety and longevity. But what factors influence this strength?

To answer this question we'll use exploratory data analysis (EDA) of concrete mix data. Through techniques like statistical summaries and visualizations, we'll unveil the distribution of each element.

Are there any outliers that stand out from the mix? We'll then delve deeper, examining how each ingredient interacts with the overall strength. Does a higher cement content necessarily translate to a stronger concrete? This initial exploration serves as a roadmap for further analysis. By understanding the data landscape, we can identify the most influential factors and potential roadblocks like missing values or inconsistencies. This knowledge will be invaluable when building models to predict concrete strength, ultimately leading to the creation of even more robust and reliable structures.

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objs as go
import statsmodels.api as sm
from scipy import stats

data = pd.read_csv('/content/drive/MyDrive/ConcreteStrengthData.csv')
data.head()

{"summary": "{\n  \"name\": \"data\",\n  \"rows\": 1030,\n  \"fields\": [\n    {\n      \"column\": \"CementComponent\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 104.50636449481532,\n        \"min\": 102.0,\n        \"max\": 540.0,\n        \"num_unique_values\": 278,\n        \"samples\": [\n          337.9,\n          290.2,\n          262.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"BlastFurnaceSlag\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 86.27934174810584,\n        \"min\": 0.0,\n        \"max\": 359.4,\n        \"num_unique_values\": 185,\n        \"samples\": [\n          94.7,\n          119.0,\n          136.3\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"FlyAshComponent\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 63.99700415268765,\n        \"min\": 0.0,\n        \"max\": 200.1,\n        \"num_unique_values\": 156,\n        \"samples\": [\n          98.0,\n          142.0,\n          195.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  }\n}
```



```

139.4\n          ],\n          \"semantic_type\": \"\",\n          \"description\": \"\"\n        },\n        {\n          \"column\": \"FlyAshComponent\",\n          \"properties\": {\n            \"dtype\": \"number\",\n            \"std\": 46.32475580075948,\n            \"min\": 0.0,\n            \"max\": 115.6,\n            \"num_unique_values\": 5,\n            \"samples\": [\n              115.6,\n              78.3,\n              108.6\n            ],\n            \"semantic_type\": \"\",\n            \"description\": \"\"\n          },\n          {\n            \"column\": \"WaterComponent\",\n            \"properties\": {\n              \"dtype\": \"number\",\n              \"std\": 10.783320453366857,\n              \"min\": 175.6,\n              \"max\": 200.6,\n              \"num_unique_values\": 5,\n              \"samples\": [\n                196.0,\n                200.6,\n                192.7\n              ],\n              \"semantic_type\": \"\",\n              \"description\": \"\"\n            },\n            {\n              \"column\": \"SuperplasticizerComponent\",\n              \"properties\": {\n                \"dtype\": \"number\",\n                \"std\": 1.9882152800941857,\n                \"min\": 6.1,\n                \"max\": 11.3,\n                \"num_unique_values\": 5,\n                \"samples\": [\n                  10.4,\n                  8.6,\n                  6.1\n                ],\n                \"semantic_type\": \"\",\n                \"description\": \"\"\n              },\n              {\n                \"column\": \"CoarseAggregateComponent\",\n                \"properties\": {\n                  \"dtype\": \"number\",\n                  \"std\": 63.48728219100265,\n                  \"min\": 817.9,\n                  \"max\": 989.6,\n                  \"num_unique_values\": 5,\n                  \"samples\": [\n                    817.9,\n                    864.5,\n                    892.4\n                  ],\n                  \"semantic_type\": \"\",\n                  \"description\": \"\"\n                },\n                {\n                  \"column\": \"FineAggregateComponent\",\n                  \"properties\": {\n                    \"dtype\": \"number\",\n                    \"std\": 20.277006682446988,\n                    \"min\": 761.5,\n                    \"max\": 813.4,\n                    \"num_unique_values\": 5,\n                    \"samples\": [\n                      813.4,\n                      761.5,\n                      780.0\n                    ],\n                    \"semantic_type\": \"\",\n                    \"description\": \"\"\n                  },\n                  {\n                    \"column\": \"AgeInDays\",\n                    \"properties\": {\n                      \"dtype\": \"number\",\n                      \"std\": 0,\n                      \"min\": 28,\n                      \"max\": 28,\n                      \"num_unique_values\": 1,\n                      \"samples\": [\n                        28\n                      ],\n                      \"semantic_type\": \"\",\n                      \"description\": \"\"\n                    },\n                    {\n                      \"column\": \"Strength\",\n                      \"properties\": {\n                        \"dtype\": \"number\",\n                        \"std\": 7.371633468913115,\n                        \"min\": 23.7,\n                        \"max\": 44.28,\n                        \"num_unique_values\": 5,\n                        \"samples\": [\n                          31.18\n                        ],\n                        \"semantic_type\": \"\",\n                        \"description\": \"\"\n                      }\n                    }\n                  }\n                }\n              }\n            }\n          }\n        }\n      ],\n      \"type\": \"dataframe\"}

```

check the size of data

```
data.shape
```

```
(1030, 9)
```

```
data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1030 entries, 0 to 1029
Data columns (total 9 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   CementComponent                      1030 non-null   float64
 1   BlastFurnaceSlag                    1030 non-null   float64
 2   FlyAshComponent                     1030 non-null   float64
 3   WaterComponent                      1030 non-null   float64
 4   SuperplasticizerComponent           1030 non-null   float64
 5   CoarseAggregateComponent            1030 non-null   float64
 6   FineAggregateComponent              1030 non-null   float64
 7   AgeInDays                          1030 non-null   int64
 8   Strength                            1030 non-null   float64
dtypes: float64(8), int64(1)
memory usage: 72.5 KB
```

Details about the data :

```
print(f"There are {data.shape[0]} instances.")
print(f"There are {data.shape[1]} dataframe columns/attributes.")
```

```
There are 1030 instances.
There are 9 dataframe columns/attributes.
```

check if we have null value

```
data.isna().sum()

CementComponent          0
BlastFurnaceSlag         0
FlyAshComponent          0
WaterComponent           0
SuperplasticizerComponent 0
CoarseAggregateComponent 0
FineAggregateComponent    0
AgeInDays               0
Strength                0
dtype: int64
```

find the duplicated

```
data.duplicated().sum()

25
```

drop duplicates

```
data.drop_duplicates()
```

```
{
  "summary": {
    "name": "data",
    "rows": 1005,
    "fields": [
      {
        "column": "CementComponent",
        "properties": {
          "dtype": "number",
          "std": 104.34426071285245,
          "min": 102.0,
          "max": 540.0,
          "num_unique_values": 278,
          "samples": [
            337.9,
            290.2,
            262.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "BlastFurnaceSlag",
        "properties": {
          "dtype": "number",
          "std": 86.17080699343222,
          "min": 0.0,
          "max": 359.4,
          "num_unique_values": 185,
          "samples": [
            94.7,
            119.0,
            136.3
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "FlyAshComponent",
        "properties": {
          "dtype": "number",
          "std": 64.20796859777064,
          "min": 0.0,
          "max": 200.1,
          "num_unique_values": 156,
          "samples": [
            98.0,
            142.0,
            195.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "WaterComponent",
        "properties": {
          "dtype": "number",
          "std": 21.339334087611302,
          "min": 121.8,
          "max": 247.0,
          "num_unique_values": 195,
          "samples": [
            195.4,
            183.8,
            127.3
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "SuperplasticizerComponent",
        "properties": {
          "dtype": "number",
          "std": 5.919966720023434,
          "min": 0.0,
          "max": 32.2,
          "num_unique_values": 111,
          "samples": [
            15.0,
            28.2,
            16.5
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "CoarseAggregateComponent",
        "properties": {
          "dtype": "number",
          "std": 77.57966673714678,
          "min": 801.0,
          "max": 1145.0,
          "num_unique_values": 284,
          "samples": [
            852.1,
            913.9,
            914.0
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "FineAggregateComponent",
        "properties": {
          "dtype": "number",
          "std": 80.34043464964552,
          "min": 594.0,
          "max": 992.6,
          "num_unique_values": 302,
          "samples": [
            710.0,
            695.4,
            769.3
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "AgeInDays",
        "properties": {
          "dtype": "number",
          "std": 63,
          "min": 1,
          "max": 365,
          "num_unique_values": 14,
          "samples": [
            91,
            100,
            28
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": ""
      }
    ]
  }
}
```



```

Ash\", \n      \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 63.99700415268765, \n          \"min\": 0.0, \n          \"max\": 200.1, \n          \"num_unique_values\": 156, \n          \"samples\": [\n          98.0, \n          142.0, \n          195.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Water\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 21.35421856503247, \n          \"min\": 121.8, \n          \"max\": 247.0, \n          \"num_unique_values\": 195, \n          \"samples\": [\n          183.8, \n          127.3\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Super plasticizer\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 5.97384139248552, \n          \"min\": 0.0, \n          \"max\": 32.2, \n          \"num_unique_values\": 111, \n          \"samples\": [\n          28.2, \n          16.5\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Coarse Aggregate\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 77.75395396672077, \n          \"min\": 801.0, \n          \"max\": 1145.0, \n          \"num_unique_values\": 284, \n          \"samples\": [\n          852.1, \n          913.9, \n          914.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Fine Aggregate\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 80.17598014240437, \n          \"min\": 594.0, \n          \"max\": 992.6, \n          \"num_unique_values\": 302, \n          \"samples\": [\n          710.0, \n          695.4, \n          769.3\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Age\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 63, \n          \"min\": 1, \n          \"max\": 365, \n          \"num_unique_values\": 14, \n          \"samples\": [\n          91, \n          100, \n          28\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      }, \n      {\n          \"column\": \"Strength\", \n          \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 16.705741961912512, \n          \"min\": 2.33, \n          \"max\": 82.6, \n          \"num_unique_values\": 845, \n          \"samples\": [\n          41.68, \n          39.59, \n          2.33\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\" \n      } \n      ] \n      }, \n      \"type\": \"dataframe\", \"variable_name\": \"data\"}

```

convert categorical variables to categories

```

data['Age'] = data['Age'].astype('category')
data.describe(include='category')

{"summary": "{\n  \"name\": \"data\", \n  \"rows\": 4, \n  \"fields\": [\n    {\n      \"column\": \"Age\", \n      \"properties\": {\n        \"dtype\": \"number\", \n        \"std\": 476, \n        \"min\": 14, \n

```

```

{"max\\": 1030,\\n          \\\"num_unique_values\\\": 4,\\n
\\\"samples\\\": [\\n          14,\\n          425,\\n          1030\\n
],\\n          \\\"semantic_type\\\": \\\"\\\",\\n          \\\"description\\\": \\\"\\\"\\n
}\\n      }\\n  ]\\n}\\",\"type\":\"dataframe\"}

data['Super plasticizer'] = data['Super
plasticizer'].astype('category')
data.describe(include='category')

{\"summary\": \"{\\n  \\\"name\\\": \\\"data\\\",\\n  \\\"rows\\\": 4,\\n  \\\"fields\\\": [\\n
    {\\n        \\\"column\\\": \\\"Super plasticizer\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n        \\\"std\\\":
461.6138357256348,\\n        \\\"min\\\": 0.0,\\n        \\\"max\\\": 1030.0,\\n
\\\"num_unique_values\\\": 4,\\n        \\\"samples\\\": [\\n        111.0,\\n
379.0,\\n        1030.0\\n        ],\\n        \\\"semantic_type\\\":
\\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n    }\\n    },\\n    {\\n
\\\"column\\\": \\\"Age\\\",\\n    \\\"properties\\\": {\\n        \\\"dtype\\\":
\\\"number\\\",\\n        \\\"std\\\": 476,\\n        \\\"min\\\": 14,\\n
\\\"max\\\": 1030,\\n        \\\"num_unique_values\\\": 4,\\n
\\\"samples\\\": [\\n        14,\\n        425,\\n        1030\\n
],\\n        \\\"semantic_type\\\": \\\"\\\",\\n        \\\"description\\\": \\\"\\\"\\n
}\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\"}

```

Top 10 Age

```

top10Age = data['Age'].value_counts().nlargest(10).to_frame()
top10Age


{\"summary\": \"{\\n  \\\"name\\\": \\\"top10Age\\\",\\n  \\\"rows\\\": 10,\\n
\\\"fields\\\": [\\n    {\\n        \\\"column\\\": \\\"Age\\\",\\n
\\\"properties\\\": {\\n        \\\"dtype\\\": \\\"category\\\",\\n
\\\"num_unique_values\\\": 10,\\n        \\\"samples\\\": [\\n        91,\\n
3,\\n        90\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n    }\\n    },\\n    {\\n        \\\"column\\\":
\\\"count\\\",\\n        \\\"properties\\\": {\\n        \\\"dtype\\\": \\\"number\\\",\\n
\\\"std\\\": 121,\\n        \\\"min\\\": 14,\\n        \\\"max\\\": 425,\\n
\\\"num_unique_values\\\": 10,\\n        \\\"samples\\\": [\\n        22,\\n
134,\\n        54\\n        ],\\n        \\\"semantic_type\\\": \\\"\\\",\\n
\\\"description\\\": \\\"\\\"\\n    }\\n    }\\n  ]\\n}\\",\"type\":\"dataframe\", \"variable_name\": \"top10Age\"}

```

Scatterplots

A scatter plot is a visual representation of how two variables relate to each other. You can use scatter plots to explore the relationship between two variables, for example by looking for any correlation between them.


```
ax = sns.scatterplot(x="Fly Ash", y="Strength", data=data)
ax.set_title("Concrete Strength vs. Fly ash")
ax.set_xlabel("Fly ash");
```



Adding color as a third dimension

we conclude that the strength of cement is the best in Age 56-120

```
sns.lmplot(x="Fly Ash", y="Strength", hue="Age", data=data);
```



```


{"properties": {"dtype": "number", "std": 0.4345399866525519, "min": -0.4, "max": 1.0, "num_unique_values": 9, "samples": [-0.15, -0.32, -0.01], "semantic_type": ""}, {"column": "Water", "properties": {"dtype": "number", "std": 0.4857840169366538, "min": -0.66, "max": 1.0, "num_unique_values": 9, "samples": [0.28, 0.11, -0.18], "semantic_type": ""}, {"column": "Super plasticizer", "properties": {"dtype": "number", "std": 0.4713396982125643, "min": -0.66, "max": 1.0, "num_unique_values": 9, "samples": [-0.19, 0.04, -0.27], "semantic_type": ""}, {"column": "Coarse Aggregate", "properties": {"dtype": "number", "std": 0.39526713892140225, "min": -0.28, "max": 1.0, "num_unique_values": 8, "samples": [-0.28, 1.0, -0.11], "semantic_type": ""}, {"column": "Fine Aggregate", "properties": {"dtype": "number", "std": 0.42839169511609865, "min": -0.45, "max": 1.0, "num_unique_values": 9, "samples": [-0.16, -0.28, -0.18], "semantic_type": ""}, {"column": "Age", "properties": {"dtype": "number", "std": 0.37625715201766524, "min": -0.19, "max": 1.0, "num_unique_values": 9, "samples": [1.0, -0.04, -0.0], "semantic_type": ""}, {"column": "Strength", "properties": {"dtype": "number", "std": 0.41517399297697394, "min": -0.29, "max": 1.0, "num_unique_values": 9, "samples": [0.33, 0.13, -0.16], "semantic_type": ""}}], "type": "dataframe"}

```

Correlation matrix to heat map

Python, and its libraries, make lots of things easy. For example, once the correlation matrix is defined, it can be passed to Seaborn's `heatmap()` method to create a heatmap (or headgrid). The basic idea of heatmaps is that they replace numbers with colors of varying shades, as indicated by the scale on the right. Cells that are lighter have higher values of r . This type of visualization can make it much easier to spot linear relationships between variables than a table of numbers. For example, if I focus on the "Strength" column, I immediately see that "Cement" and "FlyAsh" have the largest positive correlations whereas "Slag" has the large negative correlation.

```
plt.figure(figsize = (8, 10))
sns.heatmap(cormat.corr(), annot = True, fmt = '0.2f', annot_kws =
{'size' : 10}, linewidth = 2, linecolor = 'white')
plt.show()
```



Preparing the data

Define Y and X matrices

and add a constant column to the X matrix

```
import statsmodels.api as sm
Y = data['Strength']
X = data['Fly Ash']
X.head()

0    0.0
1    0.0
2    0.0
3    0.0
4    0.0
Name: Fly Ash, dtype: float64

X = sm.add_constant(X)
X.head()

{"summary": "{\n  \"name\": \"X\",\n  \"rows\": 1030,\n  \"fields\": [\n    {\n      \"column\": \"const\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.0,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Fly Ash\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 63.99700415268765,\n        \"min\": 0.0,\n        \"max\": 200.1,\n        \"num_unique_values\": 156,\n        \"samples\": [\n          98.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}", "type": "dataframe", "variable_name": "X"}
```

Running the model

```
model = sm.OLS(Y, X, missing='drop')
model_result = model.fit()
model_result.summary()

<class 'statsmodels.iolib.summary.Summary'>
"""
                                OLS Regression Results
=====
Dep. Variable:                  Strength    R-squared:
0.011
Model:                            OLS      Adj. R-squared:
0.010
Method:                    Least Squares    F-statistic:
```

```
11.63
Date: Sat, 06 Apr 2024 Prob (F-statistic):
0.000675
Time: 09:09:59 Log-Likelihood:
-4355.4
No. Observations: 1030 AIC:
8715.
Df Residuals: 1028 BIC:
8725.
Df Model: 1
```

Covariance Type: nonrobust

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          37.3139      0.679      54.978      0.000      35.982
38.646
Fly Ash        -0.0276      0.008      -3.410      0.001      -0.043
-0.012
=====
=====
```


```
Omnibus: 29.013 Durbin-Watson:
0.848
Prob(Omnibus): 0.000 Jarque-Bera (JB):
27.218
Skew: 0.351 Prob(JB):
1.23e-06
Kurtosis: 2.625 Cond. No.
110.
=====
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

"""

```
sns.histplot(model_result.resid);
```



```
sns.distplot(model_result.resid)
```

```
<ipython-input-68-73949055c77d>:1: UserWarning:
```


```
`distplot` is a deprecated function and will be removed in seaborn  
v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(model_result.resid)
```


```
<Axes: ylabel='Density'>
```



Q-Q Plot

When the quantiles of two variables are plotted against each other, then the plot obtained is known as quantile – quantile plot or qqplot. This plot provides a summary of whether the distributions of two variables are similar or not with respect to the locations.

```
sm.qqplot(model_result.resid, line='s');
```

Fit Plot

```
sm.graphics.plot_fit(model_result,1, vlines=False);
```

Fitted values versus Fly Ash

